


Training shortened response time

Matched local Qwen3-4B / MLX within each panel. Same 32 cases, sequential inference; first response excluded (31 warm).



Original adapter: 54.0% lower median and 62.2% lower p95 response time.

Panels are different experiments, not a chronological speed benchmark. Pilot adapters escalated every case and remain unapproved.